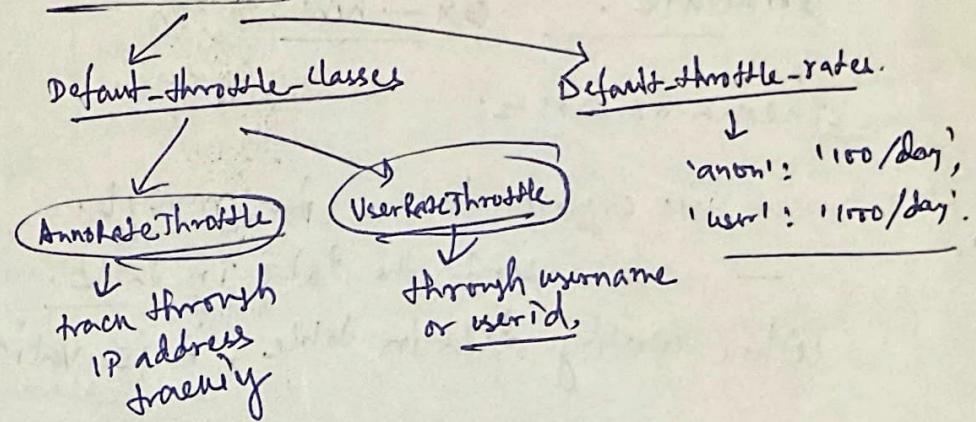


Throttling

similar to permissions.

to restrict no. of requests to your API.

two policy



Hive Indexing vs Materialize view

> Hive version 3.0

↳ there is no indexing in Hive.

to accelerate query processing. ~~in Data Warehouse~~, we should precompute relevant summaries or materialized view.

Query rewriting → it automatically rewrites incoming queries using materialize views & hence, accelerate query execution.

By default value → true.

hive.materializedview.automify in property

Maintenance of materialize view:

we can Rebuild the view when source table data changes or new data inserted.

By default Hive supports incremental rebuild & means only new data will be refreshed but to ~~go~~ refresh updated or delete operation → go with full rebuild.

* What do you mean by schema on read & schema on write? ②

schema on read :

we have the data in the form of files &
we create table structure on top of it to see the tabular
structure. Ex - hive.

schema on write:

- we create table in mysql.
- we insert the data in table.

While writing data in table, it is validating the things.

create schema in hive -

Create database mydatabase.

Comment "New Database"

Location "/path/to/db"

WITH DBPROPERTIES (property_name value).

utilizing hive partitioning with some settings. ✓

Set `hive.mapred.mode = strict`.

It restrict 3 varieties of queries -

{ ① queries on partitioned table not permitted unless you include a partition filter in WHERE clause.

② NO ORDER BY without LIMIT clause.

③ prevent cartesian product.

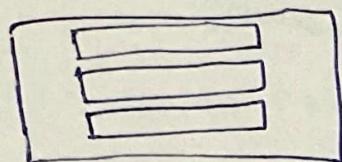
, all are permitted if you set it to "nondstrict".

How to handle data skewness in spark

⑨

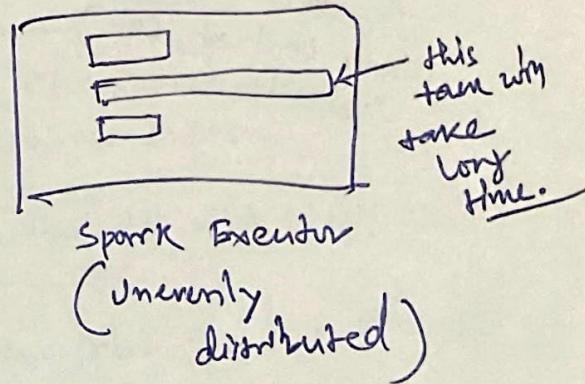
When data is unevenly distributed among partitions in the clusters.

→ it degrades performance of query, especially joins



(evenly distributed)

No of core = 3
each task one core



(Unevenly distributed)

this task will take long time.

How to handle?

- ① salt the skewed column.
- ② Apply skew-hint.
- ③ use Broadcast join.
- ④ Enable Adaptive query Execution (spark 3.0)

↓ provides three features:

- Dynamically coalescing shuffle partitions.
- Dynamically switching join strategies.
- Dynamically optimizing skew joins.

To enable -
`spark.sql.adaptive.enabled=true`

Modes in spark:

three main modes of operation:

- ① Local mode: single machine using all available core, default mode & suitable for development & testing.

- ② Standalone: when your program uses spark's resource manager, here also spark runs on cluster of machines. (for small scale production)
- ③ Cluster mode: spark runs on a cluster manager such as YARN, Mesos, etc (for large scale production).

17 SparkContext vs SparkSession:

(10)

↓
since spark 1.x

↓
you need to create
separate spark context
for every job like
live context, sql
context.

↓
since spark 2.0

↓
provides unified entry point of a
spark application.

→ all these encapsulated in
SparkSession.

Sqoop performance tuning techniques:

HDFS \longleftrightarrow RDBMS ✓

①

-num-mappers or -m (no of map tasks to run in parallel).

By default 4 mappers works.

no of map tasks < max no of connections the database supports

②

-split-by column of the table used to split work units across the mappers.

Sqoop uses this column to split the data across the executor.

If no primary key present, use -split-by argument.

③

@--boundary-query

use for creating splits.

④

-fetchsize -

No of entries to read from database at once.

By default 1000 records at a time.

⑤

-direct

fast import doesn't run any map or reduce.
only supported for mysql & postgresql.

How to handle sqoop data loss while transferring huge data?

what to do? check why failed \rightarrow get app url & see errors.

Create a staging table -

sqoop export --

--staging-table tablename

export will happen to staging table so, staging table will be imported if job failed.

--clear-staging-table \Rightarrow clear the staging table data before starting the sqoop job.

multiple small files? How to deal with it in spark?

SPM

(1) use coalesce & repartition!

optimal no of files $\rightarrow \frac{\text{Total file size}}{\text{Configured block size}}$.

(2) spark.sql.shuffle.partitions ↗

↘ set optimal value (By default 200).

Why output of the map writes to local disk, not to HDFS?

If written to HDFS \rightarrow will replicate blocks of data.

↳ overhead for namenode
to hold metadata,

↳ if job failed in between, will create
huge intermediate output residing
in HDFS & cleaning up would
be difficult,

What is speculative execution in spark.

speculative task \rightarrow is the task runs slower than the rest
of the tasks in a job,

it runs slower than median successfully
completed tasks in the task sheet.

it will not stop the slow running task but it launch new
task in parallel.

spark.speculation \rightarrow false (default)

spark.speculation.interval \rightarrow 100ms \rightarrow time interval to use before
checking for speculative tasks,

spark.speculation.multiplier \rightarrow 1.5 \rightarrow how many times slower than

spark.speculation.quantile \rightarrow .75 \rightarrow the median,

percentage of the tasks
that has not finished yet.

Where you can stored data of hive tables?

in hive metastore HDFS, S3 or any other Hadoop compatible file
system.

default \rightarrow /user/hive/warehouse.

Q1 why hdfs is not used by Hive metastore for storage?

RDBMS is being used, schema has to be updated frequently & quick access to this schema is required.

Q2 difference between external & managed tables in Hive?

↓
after drop, file will be intact.
after drop, files will also be deleted.

Q3 order by:

In strict mode, order by should provide "limit" clause.
hive.waged.mode=strict.

SORT BY:
Sort the data per reducer.

order by vs sort By

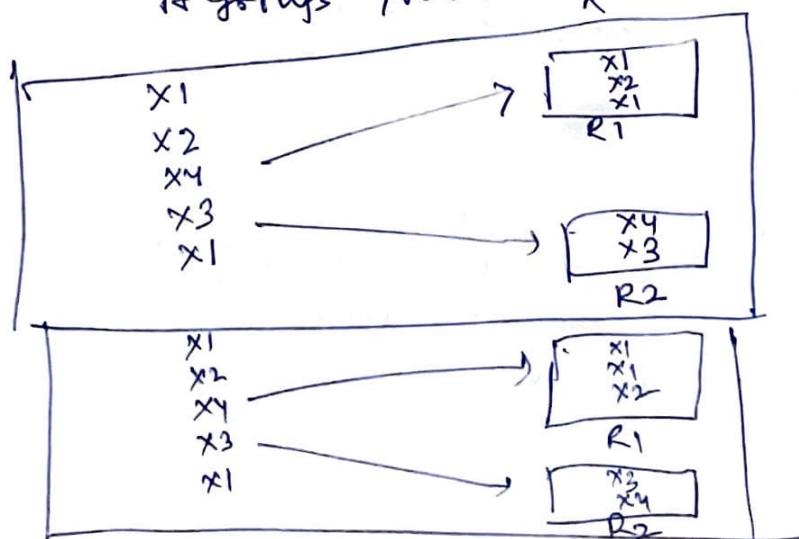
provides guarantee ordering of rows in output.

only guarantee ordering in single reducer.

if more than one reducer then Sort By gives partially ordered final result.

Distribute By:

distribute rows among reducers in a query,
it groups your data in a ^{value} single or same reducer,



Distribute By (doesn't sort in same order)

Cluster By (also does sorting)
Distribute By + Sort By

* What is dynamic partitioning in Hive?

allows for automatic creation of partition.

every time it is being taken care at runtime & no need to be aware of the data well in advance.

Steps:

1. create normal table without any partition & have the data loaded in it.
2. create a partitioned table with partition on a column name.
3. transfer the data from normal table to partitioned table

* How Bucketing works in Hive?

in Hive 3.2 → murmurhash algorithm.

you can write your custom hash fun also.

Set `hive.enforce.bucketing = true`.

Spark vs MR

100x times than MR	Faster than traditional
Scalable	Java
Batch/realtime processing	Batch processing
Caches the data in-memory & enhances the system performance	Don't support caching.

RDD

Resilient — store data lineage information to recover lost data, automatically on failure.
Distributed — RDD resides on multiple nodes. It is distributed across different nodes of a cluster.

How Spark achieves fault tolerance?
using RDD's lineage & checkpointing.

You should cache your RDD before checkpointing.

Transformation 2 actions

two types of RDD operations.

transformation — it is a function that produces new RDD from existing RDD. (map(), flatMap(), filter(),)

action — it actually works on dataset.

(Count(), collect(), save(), reduce(),).

Narrow vs wide transformation

→ no shuffling is involved (map, filter, flatmap)

shuffling is involved
reduceByKey / groupByKey

RDD vs lineage?

each RDD will have a lineage.

RDD lineage is logical execution plan & expanded every time you apply a transformation on any RDD.

DAG → combination of vertices & edges.

represents
RDDs

operations to
be applied
on RDD,

upon calling an action → DAG submitted to DAG scheduler & further splits the graph into the stages of the task.

spark partition:

to reduce data shuffle around the network in preparation for subsequent RDD processing. When the data is key-value oriented, partitioning is essential. Because the range of keys or comparable keys is in the same partition.

adaptive query execution (enabled in spark 3.2.1)

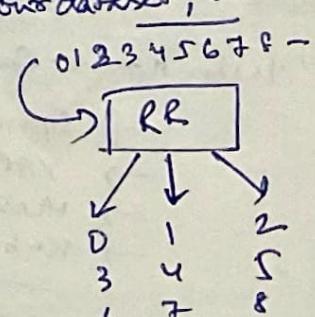
- it reduces the no of partitions used in shuffle operation in runtime
- it also determines join partition strategies in real-time & optimizes skewed join.

Still we need to do some static configuration,

we can also dynamically repartition our dataset,

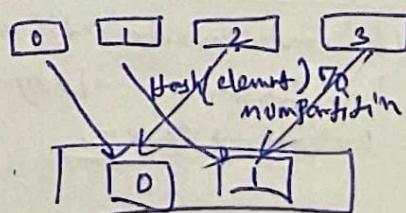
① Round Robin partitioning:

df.repartition(10)

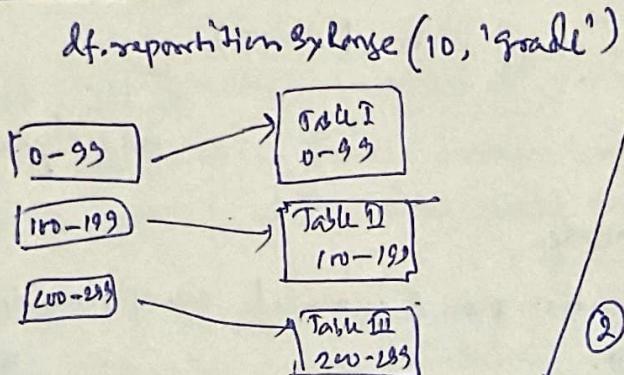


② Hash partitioning:

df.repartition(10, 'class_id')



⑤ Range partitioning:



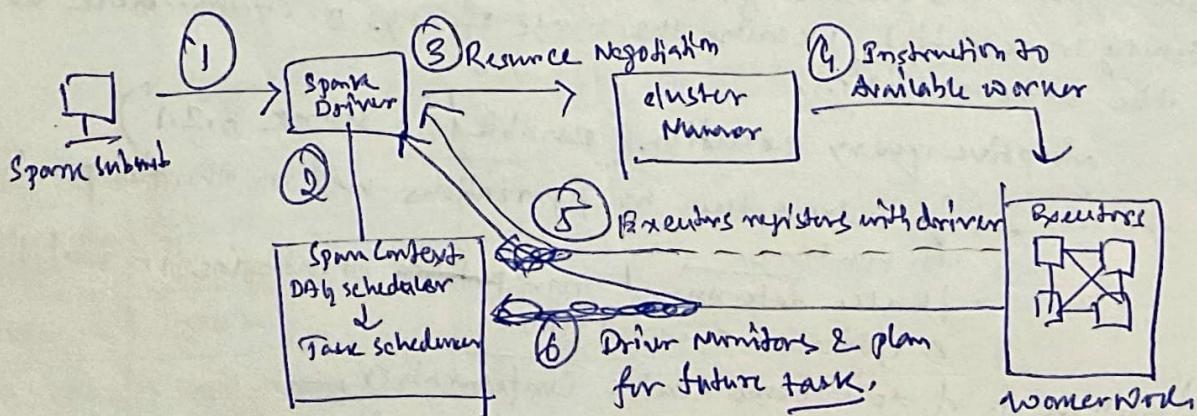
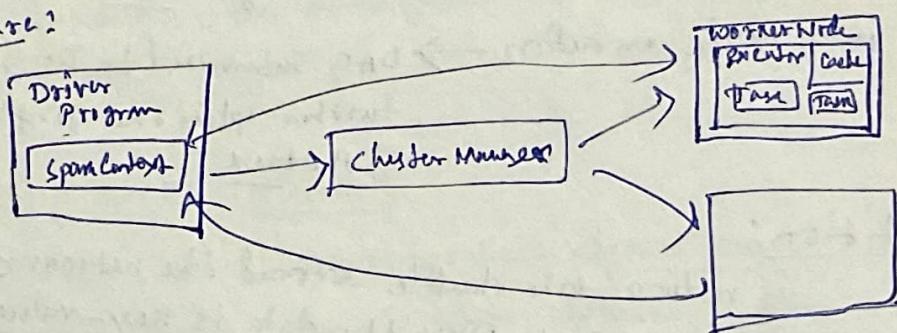
How to determine optimal no of partitions?

- ① Available resources in your cluster - ex no of available cores.
- ② How your data looks (sizing cardinality)

What is Spark Core?

distributed compute engine with all functionalities attached on its top. (Figer API lib, sparkSQL, GraphX, spark streaming).

Spark architecture:



cluster manager can be any of the following -

- spark standalone mode
- YARN
- MESOS
- Kubernetes

What is spark lazy evaluation:

spark will not execute until an ACTION is called.

transformation → Lazy
action → eager,

Diff b/w Cache & persist:

both provide an optimization mechanism to store the intermediate computation of an RDD, DF or DS, so that it can be reused.

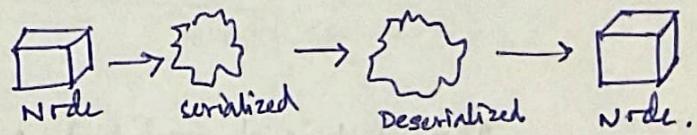
Cache()

by default saves it in memory (MEMORY-ONLY) for RDD.
(MEMORY-AND-DISK) for DF & DS.

persist()

stores in user defined storage level.

What is serialization in spark,



In distributed system, data transfer over network is the common task. If this is not handled efficiently, you may end up facing problems like, high memory usage, network bottlenecks, & performance issues.

Serialization converts objects into stream of byte in an optimal way to transfer it over nodes of network.

By default, Java serialization, but it leads to large serialization formats. We can fine-tune the performance by extending java.io.Externalizable. use Kryo serialization, takes less memory.

map vs Flatmap:

one-to-one transformation function

can return only one item

one-to-many transformation function - output is flattened.

What are the various levels of persistence in spark.

MEMORY-ONLY

MEMORY-ONLY-SER

MEMORY-AND-DISK

MEMORY-AND-DISK-SER

DISK-ONLY

OFF-HEAP

MEMORY-ONLY-2

accumulator in spark!

use a global variable in spark application. By default spark provides ~~the~~ accumulator methods for long, double & collection types.

What are broadcast variable?

→ Broadcast variable available to all executors executing the spark application.

→ size of the data you are broadcasting should be in MBs not GB.

Ex- replacing Country short form with full country name.

What is streaming checkpoint/checkpoint in spark?

- To ~~achieve~~ achieve fault tolerance in apache spark.
- It maintains RDD lineage graph & save the application state timely to reliable storage (HDFS).

Two types of checkpointing in spark

① Reliable checkpointing:

Saved to HDFS.

② Local checkpointing:

Saved to local disk in executor's memory.

RDD that is checkpointed will be computed twice, first it will compute the job, & launches another job to finish checkpoint.

So, it is recommended to do a `rdd.cache()` before `rdd.checkpoint()`, In this case 2nd job will not recompute the rdd, instead it will just read the cache.

Spark Context vs Spark Session

before spark 2.0.0,
In order to use APIs of SQL, Hive, Streaming,
separate context need to be created.

after spark 2.0.0, spark session provides a single
point of entry to interact with underlying
spark functionality.

SQL, Hive, Streaming all includes in spark session.

Memory allocation in spark

~~Explain~~
Example

Nodes \rightarrow 6

each Nodes \rightarrow 16 Cores, 64 GB RAM.

1 core & 1 GB needed for OS & Hadoop Daemons.

No of Core \Rightarrow Concurrent tasks an executor can run.
 \Rightarrow optimal value is 5.

so each node has 15 cores
63 GB RAM.

No of executors:

~~Now~~

No of executors =
(In each node)

No of cores in each Node(15)

No of cores (5)

$$= \frac{3}{5} \checkmark$$

total 6 nodes, so $6 \times 3 = 18$ executors,

- 1 executors for application master.
So, 17 executors.

- num-executors = 17.

Memory for each executor:

3 executor per node & available RAM for each node = 63 GB.

So, memory for each executor = $63 / 3 = 21$ GB.

You need small overhead memory when you request for executor memory.

that is,

$\max(384\text{MB}, 10\% \text{ of } 21\text{GB})$

$\Rightarrow 2.1$ GB.

So, executor memory to request = $21 - 2.1$ GB

≈ 18.9 GB

So, if you request 18.9 GB, then total

$(18.9 + 2.1)$ GB will be submitted to YARN. \checkmark

Stages in Spark

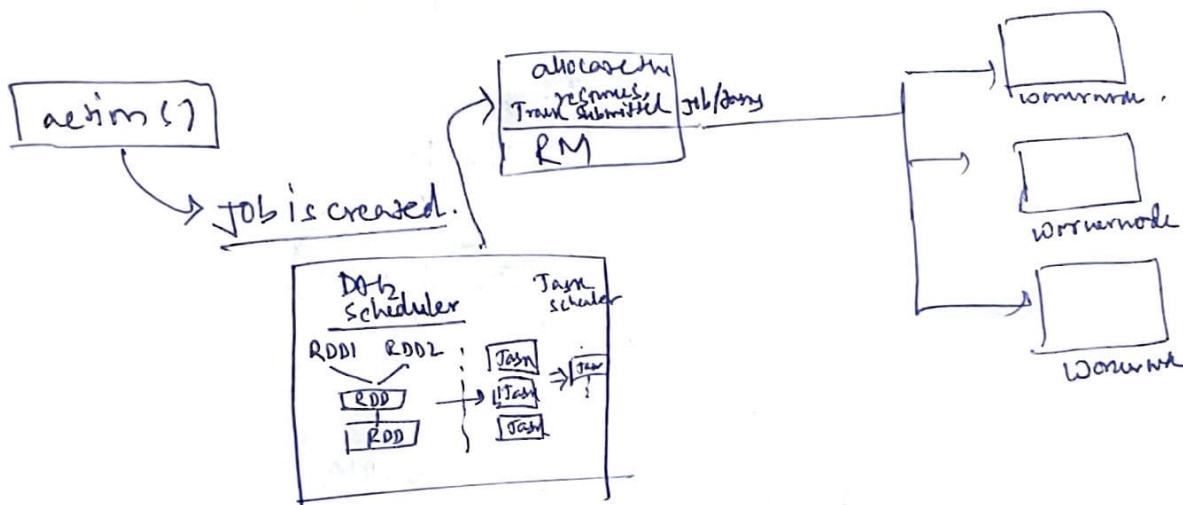
Stage is the -

- (1) physical unit of execution
- (2) step in the physical execution plan.
- (3) stage is the collection of tasks - one task per partition.
- (4) job is a collection of stages.

Two types of stages in spark -

→ Shuffle Map Stage

→ Result Stage.



⇒ We can associate stage with many other dependent parent stages.

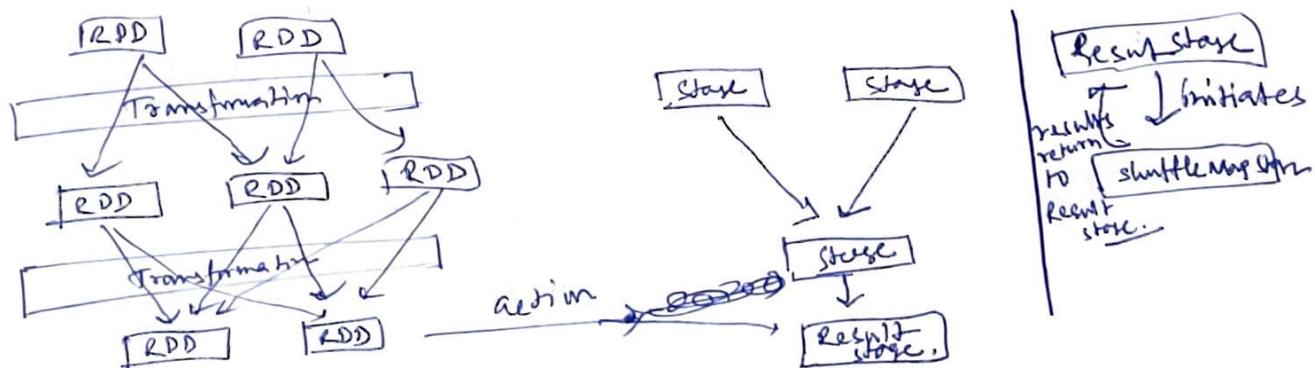
⇒ Data is passed between the stages through Exchange/shuffle.

⇒ stage submission triggers execution of a series of dependent parent stages.

Shuffle Map Stage:

Intermediate stage, produces data for another stages.

also, writes map output files for a shuffle.



Result Stage:

- A stage that executes a specific action in a user program by running a function.
- also helps for computation of the result of an action.

** Create a table structure from another table:

Create table new-table as
(select * from old-table where 1=2)

** What is a constraint & what are the seven constraint?

↳ is a set of rule which controls the data goes into the table.

Ex - primary key

Foreign key

check,

Create table students

(student-id number not null,
student-name varchar(20),
marks NUMBER check(marks>35))

Default

used to set a default value for a column.

Create table persons /

ID int not null,

Lastname varchar(255) Not null,

Age int,

city varchar(255) DEFAULT 'Kolkata'

)

Nullability:

NULL or NOTNULL.

Unique Key:

Same as primary key but Unique can accept NULL as its value.

Surrogate key:

also called synthetic primary key. It is generated when a new record is inserted into a table automatically by a database.

Ex - System date & timestamp

Random alphanumeric string.

OLTP

v/s

OLAP

online transactional processing	online analytical processing.
Highly normalized.	Highly denormalized.
<u>current data.</u>	<u>historical data.</u>
fast for delta operation.	fast for real operation.
<u>delta operations.</u>	only select operation.
	Terms Used: dimension table, fact table.

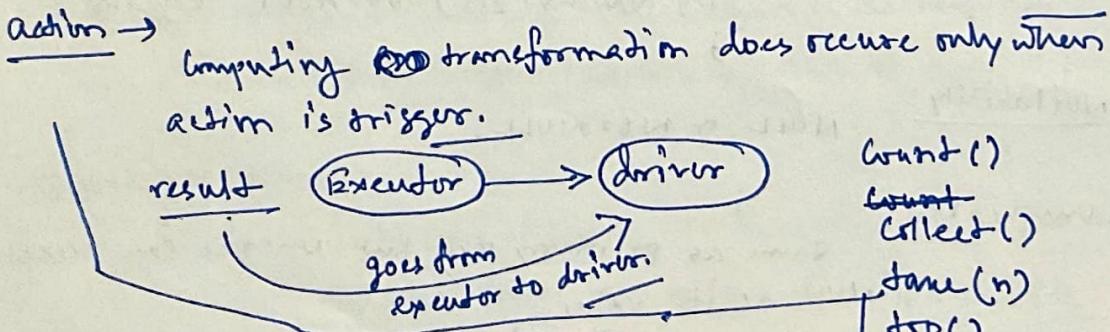
*** Explain about Cursors?

Temporary DB object which are used to loop through a table on row-by-row basis. five types of ~~cursors~~ Cursors?

*** What is action & transformation in spark?

transformation → produces new RDD from existing RDD.

- ↳ Narrow [Map, flatMap, MapPartitions, Filter, sample, Union]
- ↳ wide [Intersection, Distinct, ReduceByKey, GroupByKey, Join, Cartesian, Repartition, co-partitioned] .



count()
 count
 collect()
 take(n)
 top()
 countByValue()
 reduce()
 aggregate()
 foreach()

*** What is serialization & deserialization?

In distributed System, data transfer ~~with~~ over the network is the most common task.

If not handled properly → you may end up facing numerous problems, like ~~too~~ high memory usage, network bottlenecks,

Serialization helps to transfer data effectively over network.

→ converts object into a stream of bytes & vice versa.
to optimally transfer it over nodes of network.

Java serialization (default) / Kryo serialization (recommended by spark).

* checkpointing It is the process of saving intermediate RDDs to a reliable storage system. By periodically checkpointing RDDs, spark can recover lost data partitions & continue processing from the checkpointed state instead of recomputing from scratch.

spark

repartition & coalesce

~~val~~ $x = \text{range}(1, 13)$

$df = x.\text{toDF}(\text{"number"})$

~~val~~ $df.\text{rdd}.\text{partitions}.\text{size} \Rightarrow 4$

How data is divided into partition -

Partition 00000 : 1, 2, 3

Partition 00001 : 4, 5, 6

Partition 00002 : 7, 8, 9

Partition 00003 : 10, 11, 12.

Let's do a shuffle with repartition method & get this data on two partitions. df.repartition(2)

Partition 0 : 1, 3, 4, 6, 7, 9, 10, 12

Partition 1 : 2, 5, 8, 11.

join two data frame:

$df.\text{join}(df1, \text{on}=[\text{'session-id']}, \text{how}=\text{'inner'})$

multiple conditions.

$df.\text{join}(df2, [\text{df1.val} == \text{df2.val}, \text{df1.val2} < \text{df2.val2}], \text{how}=\text{'inner'})$

$\text{how} = \text{'left'}$
 $\text{how} = \text{'leftouter'}$
 $\text{how} = \text{'right'}$
 $\text{how} = \text{'rightouter'}$
 $\text{how} = \text{'fullouter'}$
 $\text{how} = \text{'leftsemi'}$
 $\text{how} = \text{'leftanti'}$

remove duplicates from data frame:

$df = df.\text{distinct}() \rightarrow$ drop/remove duplicate rows ~~from all columns~~ from df.

$df = df.\text{dropDuplicates}()$

or
 $df.\text{dropDuplicates}([\text{"department"}, \text{"salary"]})$

How to handle Bad or Corrupt records in Spark?

option 1:

spark.read.option("badRecordsPath", "/tmp/")

option 2:

Permissive mode

spark.read.option("mode", "PERMISSIVE")

It will process corrupted data also.
but you will see some inaccurate results
like NULL etc. for such records.

option 3:

DropMalformed mode

It excludes non-parsable records.

spark.read.option("mode", "DROPMALFORMED")

option 4:

FailFast mode:

Spark throws an exception & halts the data loading process when it finds any bad or corrupted records.

spark.read.option("mode", "FAILFAST")

option 5:

using columnNameIfCorruptRecord

spark.read.csv(columnNameIfCorruptRecord = "corrupted")

multiple delimiter

after spark 3.0

spark.read.option("delimiter", "\n").csv()

modify files in HDFS

HDFS is for write once & read multiple times.

but for append, you can use
echo "your text" | hdfs dfs -appendToFile /user/hduser/myfile.txt

-appendToFile -/user/hduser/myfile.txt

fill null values in DF:

df.na.fill(value=0)

only a particular column
df.na.fill(value=0, subset=[["particular"]])